

INDEX

S. No.	Experiment	Date	Pg No.	Sign
1	Study and Implement the Naive Bayes learner using WEKA. (Breast Cancer data file).	9/3/23	1	
2	Study and Implement the Decision Tree learners using WEKA. (Breast Cancer dataset).	16/3/23	6	
3	Estimate the accuracy of decision classifier on breast cancer dataset using 5-fold cross-validation. (You need to choose the appropriate options for missing values).	23/3/23	9	
4	Estimate the precision, recall, accuracy, and F-measure of the decision tree classifier on the text Classification task for each of the 10 categories using 10-fold cross-validation.	30/3/23	12	
5	Develop a machine learning method to classifying your incoming mail.	6/4/23	16	
6	Develop a machine learning method to Predict stock prices based on past price variation.	13/4/23	19	
7	Develop a machine learning method to Cluster gene expression data, how to modify existing methods to solve the problem better.	20/4/23	23	
8	Select two datasets. Each dataset should contain examples from multiple classes. For training purposes assume that the class label of each example is unknown (if it is known, ignore it). Implement the K-means algorithm and apply it to the data you selected. Evaluate performance by measuring the sum of Euclidean distance of each example from its class center. Test the performance of the algorithm as a function of the parameter k.	27/4/23	27	

Experiment 1

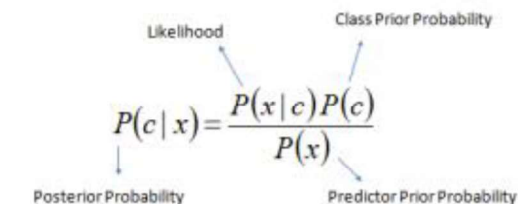
Aim: 1. Study and Implement the Naive Bayes learner using WEKA. (Breast Cancer data file).

Theory: It is a classification technique based on Bayes' Theorem with an assumption of independence among predictors. In simple terms, a Naive Bayes classifier assumes that the presence of a particular feature in a class is unrelated to the presence of any other feature.

For example, a fruit may be considered to be an apple if it is red, round, and about 3 inches in diameter. Even if these features depend on each other or upon the existence of the other features, all of these properties independently contribute to the probability that this fruit is an apple and that is why it is known as 'Naive'.

The Naive Bayes model is easy to build and particularly useful for very large data sets. Along with simplicity, Naive Bayes is known to outperform even highly sophisticated classification methods.

Bayes theorem provides a way of calculating posterior probability $P(c|x)$ from $P(c)$, $P(x)$ and $P(x|c)$. Look at the equation below:



The diagram shows the equation $P(c|x) = \frac{P(x|c)P(c)}{P(x)}$ with arrows pointing from labels to the terms: 'Likelihood' points to $P(x|c)$, 'Class Prior Probability' points to $P(c)$, 'Posterior Probability' points to $P(c|x)$, and 'Predictor Prior Probability' points to $P(x)$.

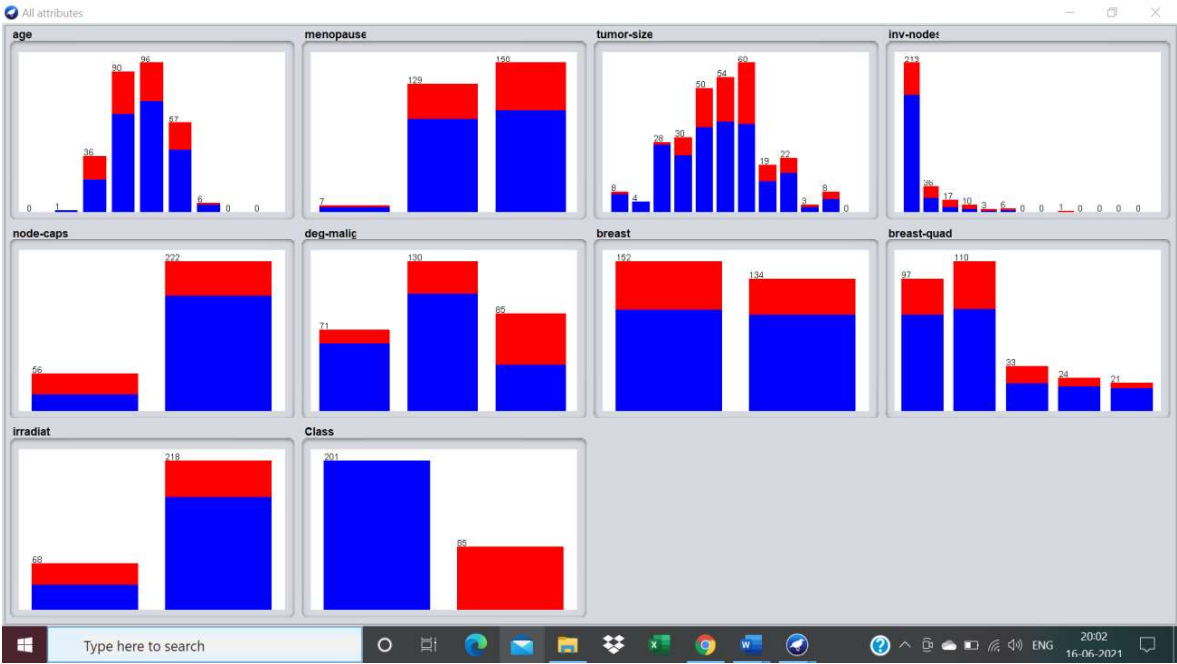
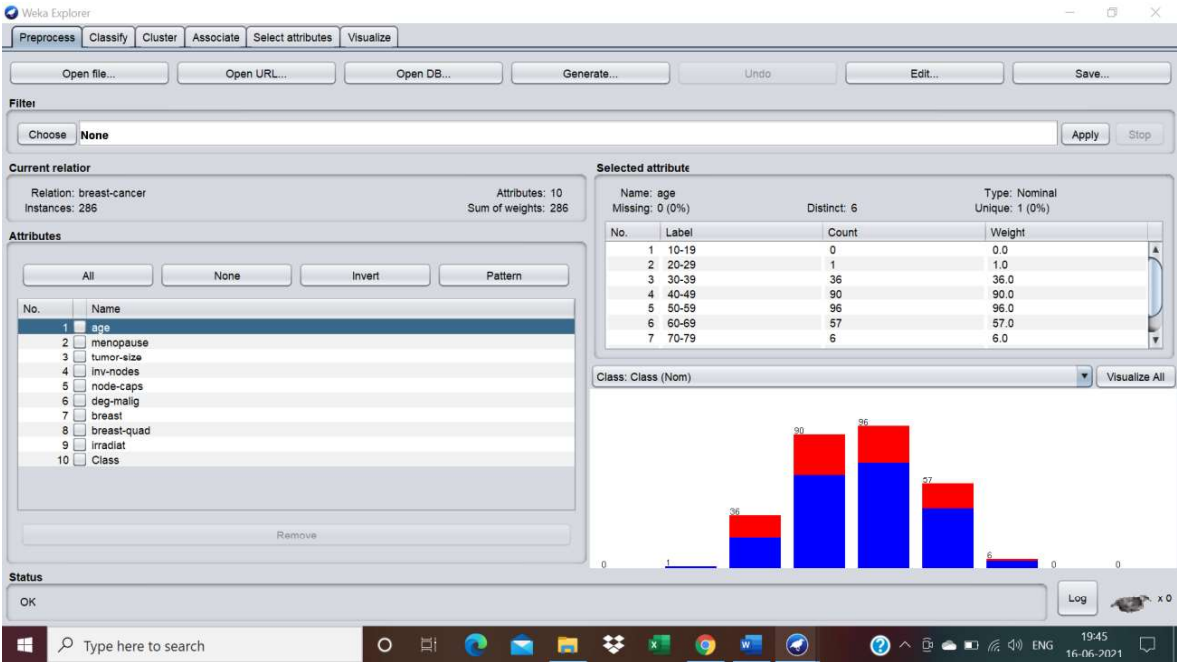
$$P(c|x) = \frac{P(x|c)P(c)}{P(x)}$$
$$P(c|X) = P(x_1|c) \times P(x_2|c) \times \dots \times P(x_n|c) \times P(c)$$

Above,

- $P(c|x)$ is the posterior probability of class (c, target) given predictor (x, attributes).
- $P(c)$ is the prior probability of class.
- $P(x|c)$ is the likelihood which is the probability of predictor given class.
- $P(x)$ is the prior probability of predictor.

First, we use the data mining tools WEKA to do the training data prediction. Here, we will use 10-fold cross validation on training data to calculate the machine learning rules and their performance. The results are as follows:

Output:



Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose NaiveBayes

Test options

☒ Use training set
☐ Supplied test set Set...
☐ Cross-validation Folds 10
☐ Percentage split % 66
 More options...

(Nom) Class

Start Stop

Result list (right-click for options)

19:46:18 - bayes.NaiveBayes

Status

OK

Log x 0

Type here to search

19:46 16-06-2021

Classifier output

```

=== Run information ===

Scheme:      weka.classifiers.bayes.NaiveBayes
Relation:    breast-cancer
Instances:   286
Attributes:  10
  age
  menopause
  tumor-size
  inv-nodes
  node-caps
  deg-malig
  breast
  breast-quad
  irradiat
  Class
Test mode:   evaluate on training data

=== Classifier model (full training set) ===

Naive Bayes Classifier

Attribute    no-recurrence-events  recurrence-events
              (0.7)          (0.3)
=====

```

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose NaiveBayes

Test options

☒ Use training set
☐ Supplied test set Set...
☐ Cross-validation Folds 10
☐ Percentage split % 66
 More options...

(Nom) Class

Start Stop

Result list (right-click for options)

19:46:18 - bayes.NaiveBayes

Status

OK

Log x 0

Type here to search

19:47 16-06-2021

Classifier output

Attribute	no-recurrence-events (0.7)	recurrence-events (0.3)
age		
10-19	1.0	1.0
20-29	2.0	1.0
30-39	22.0	16.0
40-49	64.0	28.0
50-59	72.0	26.0
60-69	41.0	18.0
70-79	6.0	2.0
80-89	1.0	1.0
90-99	1.0	1.0
[total]	210.0	94.0
menopause		
lt40	6.0	3.0
ge40	95.0	36.0
premeno	103.0	49.0
[total]	204.0	88.0
tumor-size		
0-4	8.0	2.0
5-9	5.0	1.0
10-14	28.0	2.0

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose NaiveBayes

Test options

☒ Use training set

☐ Supplied test set Set...

☐ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) Class

Start Stop

Result list (right-click for options)

19:46:18 - bayes.NaiveBayes

Classifier output

tumor-size		
0-4	8.0	2.0
5-9	5.0	1.0
10-14	28.0	2.0
15-19	24.0	8.0
20-24	35.0	17.0
25-29	37.0	19.0
30-34	36.0	26.0
35-39	13.0	8.0
40-44	17.0	7.0
45-49	3.0	2.0
50-54	6.0	4.0
55-59	1.0	1.0
[total]	213.0	97.0

inv-nodes		
0-2	168.0	47.0
3-5	20.0	16.0
6-8	8.0	11.0
9-11	5.0	7.0
12-14	2.0	3.0
15-17	4.0	4.0
18-20	1.0	1.0
21-23	1.0	1.0
24-26	1.0	2.0
27-29	1.0	1.0

Status

OK

Log x 0

Type here to search

19:48 16-06-2021

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose NaiveBayes

Test options

☒ Use training set

☐ Supplied test set Set...

☐ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) Class

Start Stop

Result list (right-click for options)

19:46:18 - bayes.NaiveBayes

Classifier output

15-17		
18-20	1.0	1.0
21-23	1.0	1.0
24-26	1.0	2.0
27-29	1.0	1.0
30-32	1.0	1.0
33-35	1.0	1.0
36-39	1.0	1.0
[total]	214.0	98.0

node-caps		
yes	26.0	32.0
no	172.0	52.0
[total]	198.0	84.0

deg-malign		
1	60.0	13.0
2	103.0	29.0
3	41.0	46.0
[total]	204.0	88.0

breast		
left	104.0	50.0
right	99.0	37.0
[total]	203.0	87.0

Status

OK

Log x 0

Type here to search

19:48 16-06-2021

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose NaiveBayes

Test options

☒ Use training set
☐ Supplied test set Set...
☐ Cross-validation Folds 10
☐ Percentage split % 66
More options...

(Nom) Class

Start Stop

Result list (right-click for options)

19:46:18 - bayes.NaiveBayes

Classifier output

```
breast
left      104.0      50.0
right     99.0      37.0
[total]   203.0      87.0

breast-quad
left_up   72.0      27.0
left_low  76.0      36.0
right_up  21.0      14.0
right_low 19.0       7.0
central   18.0       5.0
[total]  206.0      89.0

irradiat
yes       38.0      32.0
no       165.0     55.0
[total]  203.0      87.0

Time taken to build model: 0.02 seconds

=== Evaluation on training set ===

Time taken to test model on training data: 0 seconds
```

Status

OK Log x 0

Type here to search

19:49 16-06-2021

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose NaiveBayes

Test options

☒ Use training set
☐ Supplied test set Set...
☐ Cross-validation Folds 10
☐ Percentage split % 66
More options...

(Nom) Class

Start Stop

Result list (right-click for options)

19:46:18 - bayes.NaiveBayes

Classifier output

```
=== Summary ===

Correctly Classified Instances      215      75.1748 %
Incorrectly Classified Instances    71      24.8252 %
Kappa statistic                    0.3693
Mean absolute error                 0.3012
Root mean squared error             0.4278
Relative absolute error             72.0082 %
Root relative squared error        93.6095 %
Total Number of Instances          286

=== Detailed Accuracy By Class ===

      TP Rate  FP Rate  Precision  Recall  F-Measure  MDC      ROC Area  FRC Area  Class
0.866      0.518      0.798      0.866      0.831      0.374      0.760      0.879      no-recurrence-events
0.482      0.134      0.603      0.482      0.536      0.374      0.760      0.610      recurrence-events
Weighted Avg. 0.752      0.404      0.740      0.752      0.743      0.374      0.760      0.799

=== Confusion Matrix ===
  a  b  <-- classified as
174 27 | a = no-recurrence-events
 44 41 | b = recurrence-events
```

Status

OK Log x 0

Type here to search

19:46 16-06-2021

Experiment 2

Aim: Study and Implement the Decision Tree learners using WEKA. (Breast Cancer dataset).

Theory: C4.5 is an algorithm used to generate a decision tree developed by Ross Quinlan. C4.5 is an extension of Quinlan's earlier ID3 algorithm. The decision trees generated by C4.5 can be used for classification, and for this reason, C4.5 is often referred to as a statistical classifier. C4.5 builds decision trees from a set of training data in the same way as ID3, using the concept of information entropy. The training data is a set $S = \{s_1, s_2, \dots\}$ of already classified samples. Each sample s_i consists of a p -dimensional vector $\{(x_{1,i}, x_{2,i}, \dots, x_{p,i})\}$, where the $x_{j,i}$ represent attribute values or features of the sample, as well as the class in which s_i falls.

At each node of the tree, C4.5 chooses the attribute of the data that most effectively splits its set of samples into subsets enriched in one class or the other. The splitting criterion is the normalized information gain (difference in entropy). The attribute with the highest normalized information gain is chosen to make the decision. The C4.5 algorithm then recurs on the smaller sublists.

This algorithm has a few base cases.

- All the samples in the list belong to the same class. When this happens, it simply creates a leaf node for the decision tree saying to choose that class.
- None of the features provide any information gain. In this case, C4.5 creates a decision node higher up the tree using the expected value of the class.
- Instance of previously-unseen class encountered. Again, C4.5 creates a decision node higher up the tree using the expected value.

J48 is an open-source Java implementation of the C4.5 algorithm in the WEKA.

Output:

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose J48 -C 0.25 -M 2

Test options

☒ Use training set

☐ Supplied test set Set...

☐ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) Class

Start Stop

Result list (right-click for options)

20:08:53 - trees J48

Classifier output

```

=== Run information ===

Scheme:      weka.classifiers.trees.J48 -C 0.25 -M 2
Relation:    breast-cancer
Instances:   286
Attributes:  10
    age
    menopause
    tumor-size
    inv-nodes
    node-caps
    deg-malig
    breast
    breast-quad
    irradiat
    Class
Test mode:   evaluate on training data

=== Classifier model (full training set) ===

J48 pruned tree
-----

node-caps = yes
|  deg-malig = 1: recurrence-events (1.01/0.4)
|  deg-malig = 2: no-recurrence-events (26.2/8.0)
|  deg-malig = 3: recurrence-events (30.4/7.4)
node-caps = no: no-recurrence-events (228.39/53.4)

Number of Leaves :    4
Size of the tree :    6

Time taken to build model: 0.01 seconds

=== Evaluation on training set ===

Time taken to test model on training data: 0 seconds

=== Summary ===

```

Status

OK

Log x 0

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose J48 -C 0.25 -M 2

Test options

☒ Use training set

☐ Supplied test set Set...

☐ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) Class

Start Stop

Result list (right-click for options)

20:08:53 - trees J48

Classifier output

```

Test mode:   evaluate on training data

=== Classifier model (full training set) ===

J48 pruned tree
-----

node-caps = yes
|  deg-malig = 1: recurrence-events (1.01/0.4)
|  deg-malig = 2: no-recurrence-events (26.2/8.0)
|  deg-malig = 3: recurrence-events (30.4/7.4)
node-caps = no: no-recurrence-events (228.39/53.4)

Number of Leaves :    4
Size of the tree :    6

Time taken to build model: 0.01 seconds

=== Evaluation on training set ===

Time taken to test model on training data: 0 seconds

=== Summary ===

```

Status

OK

Log x 0

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose J48 - C 0.25 - M 2

Test options

☒ Use training set
☐ Supplied test set Set...
☐ Cross-validation Folds 10
☐ Percentage split % 66
More options...

(Nom) Class

Start Stop

Result list (right-click for options)

20:08:53 - trees J48

Classifier output

```

=== Summary ===
Correctly Classified Instances      217           75.8741 %
Incorrectly Classified Instances    69           24.1259 %
Kappa statistic                    0.2899
Mean absolute error                 0.3658
Root mean squared error             0.4269
Relative absolute error             87.4491 %
Root relative squared error         93.4017 %
Total Number of Instances          286

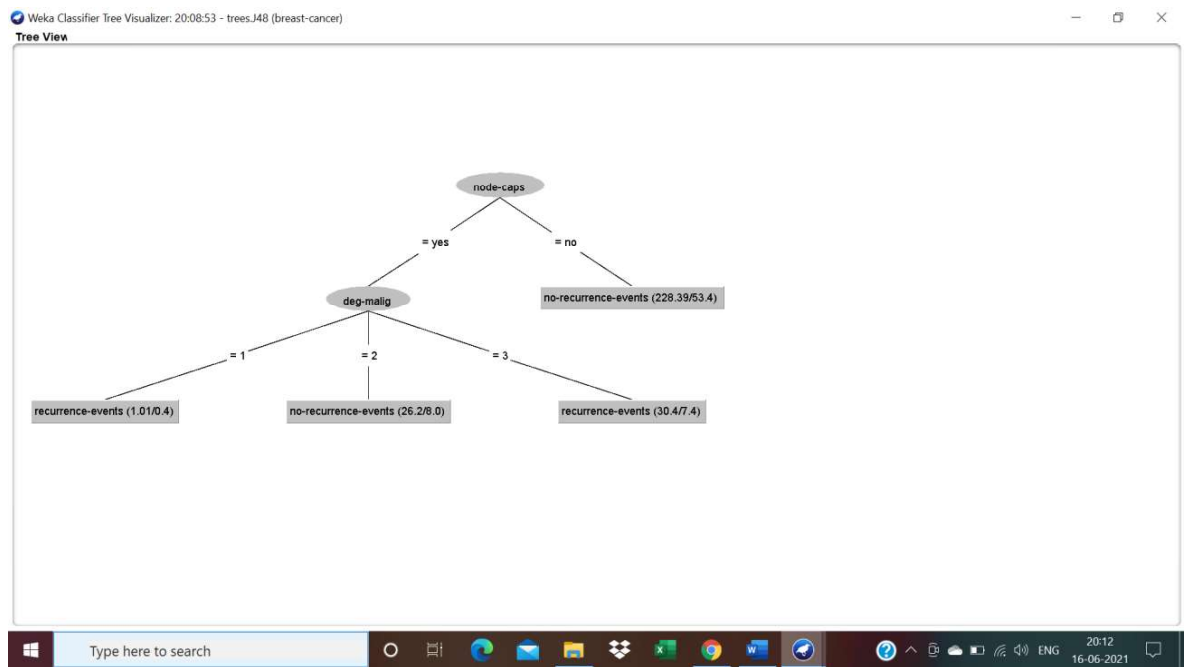
=== Detailed Accuracy By Class ===
               TP Rate  FP Rate  Precision  Recall   F-Measure  MCC     ROC Area  PRC Area  Class
0.965   0.729   0.758   0.965   0.849   0.352   0.639   0.767  no-recurrence-events
0.271   0.035   0.767   0.271   0.400   0.352   0.639   0.461  recurrence-events
Weighted Avg.   0.759   0.523   0.760   0.759   0.716   0.352   0.639   0.676

=== Confusion Matrix ===
  a  b  <-- classified as
194  7  |  a = no-recurrence-events
 62 23  |  b = recurrence-events

```

Status

OK Log x 0



Experiment 3

Aim: Estimate the accuracy of decision classifier on breast cancer dataset using 5-fold cross-validation. (You need to choose the appropriate options for missing values).

Theory: There are various approaches to determine the performance of classifiers. The performance can most simply be measured by counting the proportion of correctly predicted examples in an unseen test dataset. This value is the accuracy, which is also $1 - \text{ErrorRate}$. Both terms are used in literature. The simplest case is using a training set and a test set which are mutually independent. This is referred to as hold-out estimate. To estimate variance in these performance estimates, hold-out estimates may be computed by repeatedly re-sampling the same dataset – i.e. randomly reordering it and then splitting it into training and test sets with a specific proportion of the examples, collecting all estimates on test data and computing average and standard deviation of accuracy. A more elaborate method is cross-validation. Here, a number of folds n is specified. The dataset is randomly reordered and then split into n folds of equal size. In each iteration, one fold is used for testing and the other $n-1$ folds are used for training the classifier. The test results are collected and averaged over all folds. This gives the cross-validation estimate of the accuracy. The folds can be purely random or slightly modified to create the same class distributions in each fold as in the complete dataset. In the latter case the cross-validation is called stratified. Leave-one-out (= loo) cross-validation signifies that n is equal to the number of examples. Out of necessity, loo cross-validation has to be non-stratified, i.e. the class distributions in the test set are not related to those in the training data. Therefore loo cross-validation tends to give less reliable results. However it is still quite useful in dealing with small datasets since it utilizes the greatest amount of training data from the dataset.

Output:

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose J48 -C 0.25 -M 2

Test options

☐ Use training set

☐ Supplied test set

☒ Cross-validation Folds 5

☐ Percentage split % 66

More options...

(Nom) Class

Start Stop

Result list (right-click for options)

20.08.53 - trees.J48

20.14.22 - trees.J48

Classifier output

```
=== Stratified cross-validation ===
--- Summary ---
Correctly Classified Instances      212           74.1259 %
Incorrectly Classified Instances    74           25.8741 %
kappa statistic                    0.2288
Mean absolute error                 0.3726
Root mean squared error             0.4435
Relative absolute error             89.0412 %
Root relative squared error        97.0395 %
Total Number of Instances          286

=== Detailed Accuracy By Class ===
               TP Rate  FP Rate  Precision  Recall   F-Measure  MOC     ROC Area  PRC Area  Class
               0.960    0.776    0.745     0.960    0.839     0.287    0.582    0.728    no-recurrence-events
               0.224    0.040    0.704     0.224    0.339     0.287    0.582    0.444    recurrence-events
Weighted Avg.   0.741    0.558    0.733     0.741    0.691     0.287    0.582    0.643

=== Confusion Matrix ===
  a  b  <-- classified as
193  8  |  a = no-recurrence-events
 66 19  |  b = recurrence-events
```

Status

OK Log

Type here to search

2015 16-06-2021

Experiment 4

Aim: Estimate the precision, recall, accuracy, and F-measure of the decision tree classifier on the text Classification task for each of the 10 categories using 10-fold cross-validation.

Theory: Text classification is one of the key techniques in text mining to categorize the documents in a supervised manner. The processing of text classification involves two main problems are the extraction of feature terms that become effective keywords in the training phase and then the actual classification of the document using these feature terms in the test phase. This text classification task has numerous applications such as automated indexing of scientific articles according to predefined thesauri of technical terms, routing of customer email in a customer service department, filing patents into patent directories, automated population of hierarchical catalogues of Web resources, selective dissemination of information to consumers, identification of document genre, or detection and identification of criminal activities for military, police, or secret service environments and so on. Text classification can be used for document filtering and routing to topic-specific processing mechanisms such as information extraction and machine translation.

TP = true positives: number of examples predicted positive that are actually positive

FP = false positives: number of examples predicted positive that are actually negative

TN = true negatives: number of examples predicted negative that are actually negative

FN = false negatives: number of examples predicted negative that are actually positive

Recall is referred to as the true positive rate or sensitivity. The True Positive (TP) rate is the proportion of examples which were classified as class x, among all examples which truly have class x, i.e., how much part of the class was captured.

The Precision is the proportion of the examples which truly have class x among all those which were classified as class x.

$$\text{Recall} = \frac{tp}{tp + fn}$$

$$\text{Precision} = \frac{tp}{tp + fp}$$

$$\text{Accuracy} = \frac{tp + tn}{tp + tn + fp + fn}$$

The F-Measure is simply $2 * \text{Precision} * \text{Recall} / (\text{Precision} + \text{Recall})$, a combined measure for precision and recall.

These measures are useful for comparing classifiers.

Data file:

```
@relation textclass1
```

```
@attribute text1 {ball,goal,medals,party,poll,ministers}
```

```
@attribute text2 {wicket,ball,poll,election,performance,party}
```

```
@attribute news {politics, sports}
```

```
@data
```

```
ball, wicket, sports
```

```
goal, ball, sports
```

```
party, poll, politics
```

```
poll, election, politics
```

```
ministers, election, politics
```

```
medals, performance, sports
```

```
ball, party, sports
```

```
goal, wicket, sports
```

```
ministers, party, politics
```

```
party, election, politics
```

```
goal, election, politics
```

```
poll,performance, politics
```

```
ball,performance, sports
```

Output:

```
=== Run information ===
```

```
Scheme: weka.classifiers.trees.J48 -C 0.25 -M 2
```

```
Relation: textclass1
```

```
Instances: 13
```

```
Attributes: 3
```

text1

text2

news

Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

J48 pruned tree

text1 = ball: sports (3.0)

text1 = goal: sports (3.0/1.0)

text1 = medals: sports (1.0)

text1 = party: politics (2.0)

text1 = poll: politics (2.0)

text1 = ministers: politics (2.0)

Number of Leaves : 6

Size of the tree : 7

Time taken to build model: 0 seconds

=== Stratified cross-validation ===

=== Summary ===

<i>Correctly Classified Instances</i>	<i>4</i>	<i>30.7692 %</i>
---------------------------------------	----------	------------------

<i>Incorrectly Classified Instances</i>	<i>9</i>	<i>69.2308 %</i>
---	----------	------------------

Kappa statistic *-0.4444*
Mean absolute error *0.5192*
Root mean squared error *0.6517*
Relative absolute error *100.5319 %*
Root relative squared error *125.8013 %*
Total Number of Instances *13*

=== Detailed Accuracy By Class ===

<i>TP Rate</i>	<i>FP Rate</i>	<i>Precision</i>	<i>Recall</i>	<i>F-Measure</i>	<i>MCC</i>	<i>ROC Area</i>	<i>PRC Area</i>	<i>Class</i>
<i>0.571</i>	<i>1.000</i>	<i>0.400</i>	<i>0.571</i>	<i>0.471</i>	<i>-0.507</i>	<i>0.524</i>	<i>0.688</i>	<i>politics</i>
<i>0.000</i>	<i>0.429</i>	<i>0.000</i>	<i>0.000</i>	<i>0.000</i>	<i>-0.507</i>	<i>0.524</i>	<i>0.523</i>	<i>sports</i>

Weighted Avg. *0.308* *0.736* *0.215* *0.308* *0.253* *-0.507* *0.524* *0.612*


=== Confusion Matrix ===

a b <-- classified as

4 3 | a = politics

6 0 | b = sports

Visualize Tree



Experiment 5

Aim: Develop a machine learning method to classifying your incoming mail.

Theory: Email filtering is an extremely useful and practical problem. Email has become a mainstream form of communication. We obtain useful information each day via email but there is problem of getting unwanted emails.

Many users receive numerous unwanted emails each day which is why a spam filter has been created. The spam filter needs to sort incoming mail into wanted and unwanted. This can be tricky because the filter could allow too much spam into the inbox or could label some legitimate emails as spam. Machine learning can help solve this problem. The email client can be trained to learn where to put each email.

A machine learning algorithm for email filtering will take in a set of labeled messages as the input and will output correct labels for the testing data.

Pre-processing: Not all of the information in an email is necessary or even useful information. There is a lot of unrelated information that serves no purpose in classifying the email. This information can be eliminated from the email and not processed by the machine learning algorithm. The email header can be discarded. The user rarely sees or pays attention to this information anyhow so it can be ruled out.

Before attempting to eliminate things like stopwords or meaningless words, certain sets of data need to be preserved. Email addresses are useful and shouldn't be split apart. The data needs to be tokenized.

Words that appear with a high frequency can be ignored since they will appear in almost all samples. Common words, most punctuation, pronouns, simple verbs and adverbs can be thrown out since they will appear in almost all of the data so they can be ignored. The word "the" will appear multiple times in each sample; however, it adds no value. The words thrown out make up a stoplist which there are many predefined lists of stopwords that can be used.

After stop words are discarded, the remaining words are stemmed. Stemming reduces words to their root word. This allows for better comparisons between words. If the words "running" and "runner" both appeared, they would be treated as two different features. By stemming, they are both reduced to "run" which allows them to be grouped together. Stemming helps the program to correctly classify an email. The algorithm must depend on multiple words to create a classification for the email.

Techniques: There are many machine learning techniques available to filter emails. Many spam filtering techniques use text categorization methods to label the data. Using machine learning can easily be applied to this problem.

This section will give a brief overview of each algorithm how it can be used for this machine learning problem. The algorithm that will be discussed is Naive Bayes.

Data used:

	the	to	ect	and	for	of	a	you	hou	in	...	connevey	jay	valued	lay	infrastructure	military	allowing	ff	dry	Prediction
0	0	0	1	0	0	0	2	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	8	13	24	6	6	2	102	1	27	18	...	0	0	0	0	0	0	0	1	0	0
2	0	0	1	0	0	0	8	0	0	4	...	0	0	0	0	0	0	0	0	0	0
3	0	5	22	0	5	1	51	2	10	1	...	0	0	0	0	0	0	0	0	0	0
4	7	6	17	1	5	2	57	0	9	3	...	0	0	0	0	0	0	0	1	0	0

Source code:

```
import os
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.naive_bayes import GaussianNB
```

```
from sklearn.svm import SVC
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
data_dir = 'data'
```

```
filename = 'emails.csv'
```

```
filepath = os.path.join(data_dir, filename)
```

```
emails = pd.read_csv(filepath)
```

```
emails.drop(columns='Email No.', inplace=True)
```

```
print('Email dataset with {} rows and {} columns'.format(*emails.shape))
```

```
# emails.head()
```

```
X = emails.iloc[:, :-1]
y = emails.iloc[:, -1].values
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.8, random_state=42)
print(f'Training set consists of {X_train.shape[0]} rows of data')
print(f'Test set consists of {X_test.shape[0]} rows of data')
```

```
## Naive Bayes
```

```
model = GaussianNB()
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

```
acc = accuracy_score(y_test, y_pred)
print(f'Accuracy: {acc:%}')
```

Output:

```
(msit) PS C:\Users\kartik-mehra\Documents\sem-8\ML> python .\spam_classifier.py
Email dataset with 5172 rows and 3001 columns
Training set consists of 4137 rows of data
Test set consists of 1035 rows of data
Accuracy: 94.782609%
```

Experiment 6

Aim: Develop a machine learning method to Predict stock prices based on past price variation.

Theory: Stock market prediction is the act of trying to determine the future value of a company stock or other financial instrument traded on an exchange. The successful prediction of a stock's future price could yield significant profit. The efficient-market hypothesis suggests that stock prices reflect all currently available information and any price changes that are not based on newly revealed information thus are inherently unpredictable. Others disagree and those with this viewpoint possess myriad methods and technologies which purportedly allow them to gain future price information.

Data used:

	Date	Open	High	Low	Last	Close	Total Trade Quantity	Turnover (Lacs)
0	2018-10-08	208.00	222.25	206.85	216.00	215.15	4642146.0	10062.83
1	2018-10-05	217.00	218.60	205.90	210.25	209.20	3519515.0	7407.06
2	2018-10-04	223.50	227.80	216.15	217.25	218.20	1728786.0	3815.79
3	2018-10-03	230.00	237.50	225.75	226.45	227.60	1708590.0	3960.27
4	2018-10-01	234.55	234.60	221.05	230.30	230.90	1534749.0	3486.05

Source code:

```
import os

os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error

import tensorflow as tf
```

```
data_dir = 'data'
```

```
df = pd.read_csv(os.path.join(data_dir, 'NSE-TATAGLOBAL11.csv'), parse_dates=['Date'],  
index_col='Date', usecols=['Date', 'Close'])
```

```
df.sort_index(inplace=True)
```

```
days = 100
```

```
train_size = int(0.8 * df.shape[0])
```

```
scaler = MinMaxScaler()
```

```
dataset = scaler.fit_transform(df.values)
```

```
train = dataset[:train_size]
```

```
test = dataset[train_size - days:]
```

```
X_train, y_train = [], []
```

```
for i in range(days, train.shape[0]):
```

```
    X_train.append(train[i - days : i])
```

```
    y_train.append(train[i, 0])
```

```
X_train = np.array(X_train)
```

```
y_train = np.array(y_train)
```

```
X_test, y_test = [], []
```

```
for i in range(days, test.shape[0]):
```

```
    X_test.append(test[i - days : i])
```

```
    y_test.append(test[i, 0])
```

```
X_test = np.array(X_test)
```

```
y_test = np.array(y_test)
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.LSTM(64, return_sequences=True, input_shape=(days, 1)),
    tf.keras.layers.LSTM(64),
    tf.keras.layers.Dense(1)
])

model.compile(
    optimizer='adam',
    loss='mse'
)

model.fit(X_train, y_train, epochs=1)

y_pred = model.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f'Root mean squared error on test data: {rmse}')


# plt.figure(figsize=(16, 8))
plt.plot(pd.DataFrame(data=scaler.inverse_transform(train), index=df.index[:train_size]))
plt.plot(pd.DataFrame(data=scaler.inverse_transform(test[days:]),
index=df.index[train_size:]))
plt.plot(pd.DataFrame(data=scaler.inverse_transform(y_pred), index=df.index[train_size:]))

plt.xlabel('Time')
plt.ylabel('Price')
plt.title('Stock price prediction')

plt.show()
```

Output:

```
(msit) PS C:\Users\kartik-mehra\Documents\sem-8\ML> python .\stock_prediction.py  
28/28 [=====] - 4s 132ms/step - loss: 0.0071  
Root mean squared error on test data: 0.09530719443253376
```



Experiment 7

Aim: Develop a machine learning method to Cluster gene expression data, how to modify existing methods to solve the problem better.

Theory: k-nearest neighbors (kNN) are a classical method for recommender systems. Ratings or items are predicted by using the past ratings/items of the k most similar users and/or items. If the influence of the similar users/items is weighted by the similarity, all users/items may be used for the prediction. Popular similarity metrics are the Pearson correlation and the cosine similarity. Using the user-item matrix to compute the similarity is often called collaborative filtering. Computing the item similarities from the item attributes leads to content-based filtering.

The k nearest neighbors (k-NN) algorithm uses an idea distance between each instance. The distance between two messages can be measured and it can be determined how close they are to each other. The algorithm attempts to classify the message by looking at the nearest neighbors.

The distance is usually measured by using the Euclidian distance. The formula for the Euclidian distance is given by

$$d(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2}.$$

The goal is to measure the distance of k closest neighbors and take the majority vote of those neighbors to classify the current message.

Training: To create the training data simply store the feature vectors and the label of each message.

Classification: Take a message m, find the k nearest neighbors and count the number of each label (spam or not) that are given from the neighbors. If there are more spam messages in the k nearest neighbors then it is classified as spam. If not, then that message is classified as legitimate mail.

One advantage of this algorithm is that there isn't really a training phase. However, to classify a message, all distances between that message and all the training examples must be calculated and the k nearest neighbors need to be found and counted.

Dataset: MovieLens

Source code:

```
import csv  
  
import random  
  
import math
```



```
import operator
```

```
def loadDataset(filename, split, trainingSet=[], testSet=[]):
```

```
    with open(filename, 'rb') as csvfile:
```

```
        lines = csv.reader(csvfile)
```

```
        dataset = list(lines)
```

```
        for x in range(len(dataset)-1):
```

```
            for y in range(4):
```

```
                dataset[x][y] = float(dataset[x][y])
```

```
            if random.random() < split:
```

```
                trainingSet.append(dataset[x])
```

```
            else:
```

```
                testSet.append(dataset[x])
```

```
def euclideanDistance(instance1, instance2, length):
```

```
    distance = 0
```

```
    for x in range(length):
```

```
        distance += pow((instance1[x] - instance2[x]), 2)
```

```
    return math.sqrt(distance)
```

```
def getNeighbors(trainingSet, testInstance, k):
```

```
    distances = []
```

```
    length = len(testInstance)-1
```

```
    for x in range(len(trainingSet)):
```

```
        dist = euclideanDistance(testInstance, trainingSet[x], length)
```

```
        distances.append((trainingSet[x], dist))
```

```
        distances.sort(key=operator.itemgetter(1))
```

```
    neighbors = []
```

```

        for x in range(k):
            neighbors.append(distances[x][0])
    return neighbors

def getResponse(neighbors):
    classVotes = {}
    for x in range(len(neighbors)):
        response = neighbors[x][-1]
        if response in classVotes:
            classVotes[response] += 1
        else:
            classVotes[response] = 1
    sortedVotes = sorted(classVotes.iteritems(), key=operator.itemgetter(1), reverse=True)
    return sortedVotes[0][0]

def getAccuracy(testSet, predictions):
    correct = 0
    for x in range(len(testSet)):
        if testSet[x][-1] is predictions[x]:
            correct += 1
    return (correct/float(len(testSet))) * 100.0

# prepare data
trainingSet=[]
testSet=[]
split = 0.67
loadDataset('iris.data', split, trainingSet, testSet)
print 'Train set: ' + repr(len(trainingSet))
print 'Test set: ' + repr(len(testSet))

```

```
# generate predictions
predictions=[]
k = 3
for x in range(len(testSet)):
    neighbors = getNeighbors(trainingSet, testSet[x], k)
    result = getResponse(neighbors)
    predictions.append(result)
print('> predicted=' + repr(result) + ', actual=' + repr(testSet[x][-1]))
accuracy = getAccuracy(testSet, predictions)
print('Accuracy: ' + repr(accuracy) + '%')
```

Experiment 8

Aim: Select two datasets. Each dataset should contain examples from multiple classes. For training purposes assume that the class label of each example is unknown (if it is known, ignore it). Implement the K-means algorithm and apply it to the data you selected. Evaluate performance by measuring the sum of Euclidean distance of each example from its class center. Test the performance of the algorithm as a function of the parameter k .

Theory: Clustering is one of the most common exploratory data analysis technique used to get an intuition about the structure of the data. It can be defined as the task of identifying subgroups in the data such that data points in the same subgroup (cluster) are very similar while data points in different clusters are very different. In other words, we try to find homogeneous subgroups within the data such that data points in each cluster are as similar as possible according to a similarity measure such as euclidean-based distance or correlation-based distance. The decision of which similarity measure to use is application-specific.

Clustering analysis can be done on the basis of features where we try to find subgroups of samples based on features or on the basis of samples where we try to find subgroups of features based on samples. We'll cover here clustering based on features. Clustering is used in market segmentation; where we try to find customers that are similar to each other whether in terms of behaviors or attributes, image segmentation/compression; where we try to group similar regions together, document clustering based on topics, etc.

Unlike supervised learning, clustering is considered an unsupervised learning method since we don't have the ground truth to compare the output of the clustering algorithm to the true labels to evaluate its performance. We only want to try to investigate the structure of the data by grouping the data points into distinct subgroups.

k-means algorithm is an iterative algorithm that tries to partition the dataset into K pre-defined distinct non-overlapping subgroups (clusters) where each data point belongs to only one group. It tries to make the intra-cluster data points as similar as possible while also keeping the clusters as different (far) as possible. It assigns data points to a cluster such that the sum of the squared distance between the data points and the cluster's centroid (arithmetic mean of all the data points that belong to that cluster) is at the minimum. The less variation we have within clusters, the more homogeneous (similar) the data points are within the same cluster.

The way k-means algorithm works is as follows:

1. Specify number of clusters K .
2. Initialize centroids by first shuffling the dataset and then randomly selecting K data points for the centroids without replacement.

3. Keep iterating until there is no change to the centroids. i.e assignment of data points to clusters isn't changing.
 - Compute the sum of the squared distance between data points and all centroids.
 - Assign each data point to the closest cluster (centroid).
 - Compute the centroids for the clusters by taking the average of the all the data points that belong to each cluster.

Source code:

```
import numpy as np

import matplotlib.pyplot as plt

from sklearn.cluster import KMeans

from sklearn.datasets import make_blobs


def measureKMeansPerformance(X, n_samples, n_features, centers, max_clusters=10,
random_state=42):

    sum_squared_distances = []

    euclidean_distances = []

    clusters = []


    for n_clusters in range(1, max_clusters + 1):

        model = KMeans(n_clusters=n_clusters, random_state=random_state)

        model.fit(X)

        assigned_centers = model.cluster_centers_[model.predict(X)]

        euclidean_distance = np.sqrt(np.square(assigned_centers - X).sum(axis=1)).sum()

        euclidean_distances.append(euclidean_distance)

        sum_squared_distances.append(model.inertia_)

        clusters.append(n_clusters)


plt.plot(clusters, euclidean_distances)

plt.xlabel('Number of clusters')
```

```

plt.ylabel('Euclidean distance')

plt.title(f'KMeans clustering with {centers} centers')

plt.show()

return {cluster: euclidean_distance for cluster, euclidean_distance in zip(clusters,
euclidean_distances)}

row_format = "{:<15}" * 2

# Dataset 1
n_samples = 10000
n_features = 2
centers = 7

X1, _ = make_blobs(n_samples=n_samples, n_features=n_features, centers=centers,
random_state=42)

result1 = measureKMeansPerformance(X1, n_samples, n_features, centers,
max_clusters=15)

print('Result of dataset 1')
print('-'*30)
print(row_format.format('Cluster', 'Distance'))
for item in result1.items():
    print(row_format.format(*item))

# Dataset 2
n_samples = 30000
n_features = 3
centers = 4

```

```

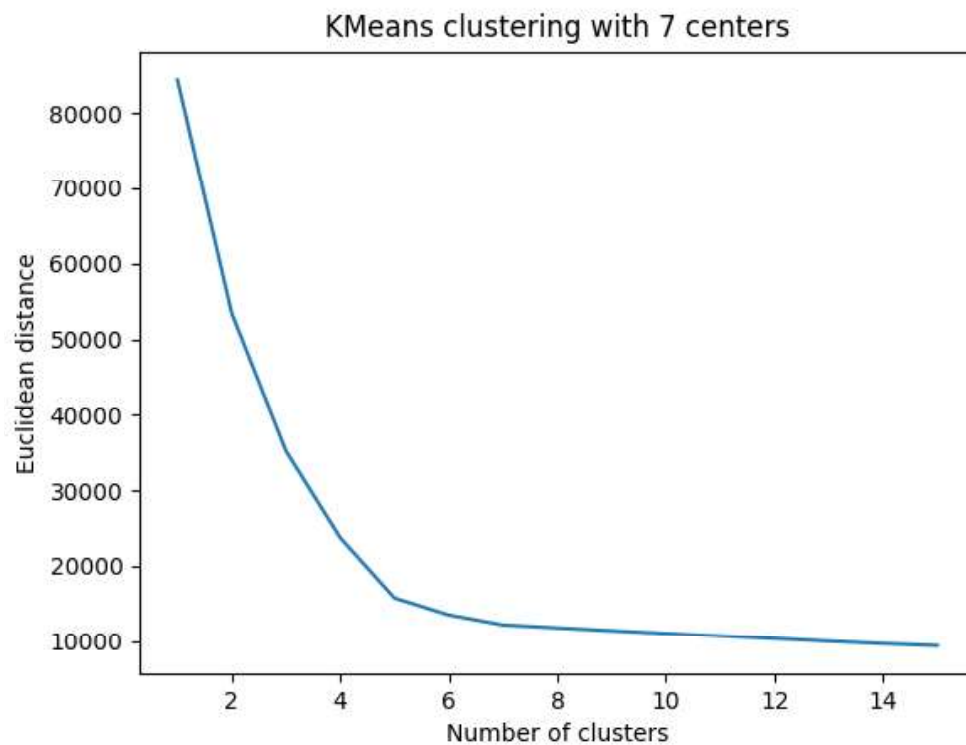
X2, _ = make_blobs(n_samples=n_samples, n_features=n_features, centers=centers,
random_state=42)

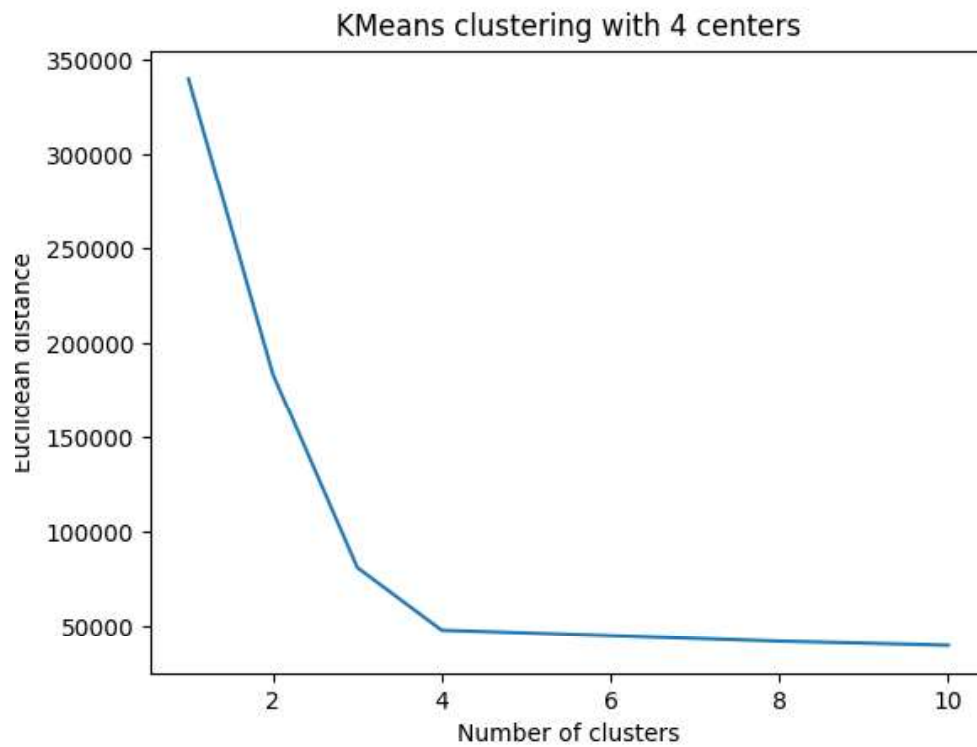
result2 = measureKMeansPerformance(X2, n_samples, n_features, centers)

print()
print('*' * 30, end='\n\n')
print('Result of dataset 2')
print('-'*30)
print(row_format.format('Cluster', 'Distance'))
for item in result2.items():
    print(row_format.format(*item))

```

Output:





```
(msit) PS C:\Users\kartik-mehra\Documents\sem-8\ML> python .\kmeans.py
Result of dataset 1
```

```
-----
Cluster      Distance
1            84385.61829910308
2            53422.63380084099
4            23749.149644736157
5            15773.49558747096
6            13522.001028897645
7            12198.828231856707
9            11422.821714504087
10           11056.427653565215
11           10705.131790052816
12           10353.990155362075
13           10000.027031073116
14           9684.676500082958
15           9414.639175310775
```

```
*****
```

```
Result of dataset 2
```

```
-----
Cluster      Distance
1            340061.4239153055
2            183791.26628029608
3            81143.63893188817
4            47959.85551576039
6            45038.13557302422
7            43608.00123766728
8            42130.06231067693
9            41082.093219088165
10           39973.21077156419
```